

Министерство образования Республики Беларусь  
Учреждение образования  
«Брестский Государственный технический университет»  
Кафедра ИИТ

Лабораторная работа №1

По дисциплине «ОИвИС»

Тема: “Обучение классификаторов средствами библиотеки PyTorch”

Выполнил:  
Студент 4 курса  
Группы ИИ-21  
Кирилович А. А.  
Проверил:  
Крощенко А.А.

Брест 2024

Вариант 10.

Выборка: STL-10 (размеченная часть).

Размер исходного изображения: 96\*96

Оптимизатор: Adam.

Цель: научиться конструировать нейросетевые классификаторы и выполнять их обучение на известных выборках компьютерного зрения

Задание 1. Выполнить конструирование своей модели СНС, обучить ее на выборке по заданию (использовать **torchvision.datasets**). Предпочтение отдавать как можно более простым архитектурам, базирующимся на базовых типах слоев (сверточный, полносвязный, подвыборочный, слой нелинейного преобразования). Оценить эффективность обучения на тестовой выборке, построить график изменения ошибки (matplotlib);

Задание 2. Реализовать визуализацию работы СНС из пункта 1 (выбор и подачу на архитектуру произвольного изображения с выводом результата);

In [ ]:

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import numpy as np
```

In [113]:

```
from torchvision import transforms

train_transform = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(10),
    transforms.RandomResizedCrop(96, scale=(0.8, 1.0)),
    transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

train_loader = torch.utils.data.DataLoader(
    torchvision.datasets.STL10('/files/', split='train', folds=1, download=True,
                               transform=train_transform),
    batch_size=64, shuffle=True)

test_loader = torch.utils.data.DataLoader(
    torchvision.datasets.STL10('/files/', split='test', folds=1, download=True,
                               transform=test_transform),
    batch_size=64, shuffle=True)
```

Files already downloaded and verified

Files already downloaded and verified

In [114]:

```
class NN(nn.Module):
    def __init__(self):
        super(NN, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        self.bn1 = nn.BatchNorm2d(32)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.bn2 = nn.BatchNorm2d(64)
        self.pool = nn.MaxPool2d(2, 2)
        # self.dropout1 = nn.Dropout(0.1)

        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1)
        self.bn3 = nn.BatchNorm2d(128)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1)
        self.bn4 = nn.BatchNorm2d(256)
        # self.dropout2 = nn.Dropout(0.2)

        self.conv5 = nn.Conv2d(256, 512, kernel_size=3, stride=1, padding=1)
        self.bn5 = nn.BatchNorm2d(512)
        self.conv6 = nn.Conv2d(512, 512, kernel_size=3, stride=1, padding=1)
        self.bn6 = nn.BatchNorm2d(512)
        # self.dropout3 = nn.Dropout(0.1)

        self.global_avg_pool = nn.AdaptiveAvgPool2d((1, 1))
```

```

self.fc1 = nn.Linear(512, 256)
self.fc2 = nn.Linear(256, 128)
self.fc3 = nn.Linear(128, 10)
self.relu = nn.LeakyReLU()
self.dropout4 = nn.Dropout(0.5)

def forward(self, x):
    x = self.pool(self.relu(self.bn1(self.conv1(x))))
    x = self.pool(self.relu(self.bn2(self.conv2(x))))
    # x = self.dropout1(x)

    x = self.pool(self.relu(self.bn3(self.conv3(x))))
    x = self.pool(self.relu(self.bn4(self.conv4(x))))
    # x = self.dropout2(x)

    x = self.pool(self.relu(self.bn5(self.conv5(x))))
    x = self.pool(self.relu(self.bn6(self.conv6(x))))
    # x = self.dropout3(x)

    x = self.global_avg_pool(x)
    x = x.view(-1, 512)

    x = self.relu(self.fc1(x))
    x = self.dropout4(x)
    x = self.relu(self.fc2(x))
    x = self.fc3(x)

    return x

```

```
model = NN()
```

In [115]:

```

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.00003)

```

In [116]:

```

def train(model, loader, criterion, optimizer, device):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    for images, labels in loader:
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()

        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

        _, predicted = torch.max(outputs, 1)
        correct += (predicted == labels).sum().item()
        total += labels.size(0)

    accuracy = 100 * correct / total
    return running_loss / len(loader), accuracy

def test(model, loader, criterion, device):
    model.eval()
    running_loss = 0.0
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in loader:

```

```

images, labels = images.to(device), labels.to(device)
outputs = model(images)
loss = criterion(outputs, labels)
running_loss += loss.item()

_, predicted = torch.max(outputs, 1)
correct += (predicted == labels).sum().item()
total += labels.size(0)

accuracy = 100 * correct / total
return running_loss / len(loader), accuracy

```

In [117]:

```

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = model.to(device)

train_losses = []
test_losses = []
train_accuracies = []
test_accuracies = []

num_epochs = 100

for epoch in range(num_epochs):
    train_loss, train_accuracy = train(model, train_loader, criterion, optimizer, device)
    test_loss, test_accuracy = test(model, test_loader, criterion, device)

    train_losses.append(train_loss)
    test_losses.append(test_loss)
    train_accuracies.append(train_accuracy)
    test_accuracies.append(test_accuracy)

    print(f'Epoch {epoch+1}/{num_epochs}, Train Loss: {train_loss:.4f}, Train Accuracy: {train_accuracy:.2f}%', '
          f'Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy:.2f}%')

```

```

Epoch 1/100, Train Loss: 2.2743, Train Accuracy: 13.70%, Test Loss: 2.3016, Test Accuracy
: 10.04%
Epoch 2/100, Train Loss: 2.2081, Train Accuracy: 21.00%, Test Loss: 2.2265, Test Accuracy
: 23.31%
Epoch 3/100, Train Loss: 2.1602, Train Accuracy: 22.50%, Test Loss: 2.0950, Test Accuracy
: 28.14%
Epoch 4/100, Train Loss: 2.0936, Train Accuracy: 24.50%, Test Loss: 2.0226, Test Accuracy
: 28.80%
Epoch 5/100, Train Loss: 2.0355, Train Accuracy: 27.00%, Test Loss: 1.9462, Test Accuracy
: 34.01%
Epoch 6/100, Train Loss: 1.9910, Train Accuracy: 27.90%, Test Loss: 1.8786, Test Accuracy
: 35.24%
Epoch 7/100, Train Loss: 1.9190, Train Accuracy: 34.30%, Test Loss: 1.8194, Test Accuracy
: 36.14%
Epoch 8/100, Train Loss: 1.8716, Train Accuracy: 31.90%, Test Loss: 1.7578, Test Accuracy
: 39.09%
Epoch 9/100, Train Loss: 1.8236, Train Accuracy: 35.10%, Test Loss: 1.6979, Test Accuracy
: 41.23%
Epoch 10/100, Train Loss: 1.7695, Train Accuracy: 37.10%, Test Loss: 1.6722, Test Accurac
y: 40.79%
Epoch 11/100, Train Loss: 1.7012, Train Accuracy: 39.60%, Test Loss: 1.6099, Test Accurac
y: 42.55%
Epoch 12/100, Train Loss: 1.6775, Train Accuracy: 38.90%, Test Loss: 1.5788, Test Accurac
y: 45.09%
Epoch 13/100, Train Loss: 1.6165, Train Accuracy: 42.50%, Test Loss: 1.5522, Test Accurac
y: 45.66%
Epoch 14/100, Train Loss: 1.6017, Train Accuracy: 41.30%, Test Loss: 1.5178, Test Accurac
y: 45.74%
Epoch 15/100, Train Loss: 1.5592, Train Accuracy: 43.00%, Test Loss: 1.5186, Test Accurac
y: 44.77%
Epoch 16/100, Train Loss: 1.5298, Train Accuracy: 45.10%, Test Loss: 1.5113, Test Accurac
y: 46.21%
Epoch 17/100, Train Loss: 1.5019, Train Accuracy: 45.60%, Test Loss: 1.4579, Test Accurac
y: 47.69%

```

Epoch 18/100,	Train Loss: 1.4765,	Train Accuracy: 47.30%,	Test Loss: 1.4398,	Test Accurac
y: 48.14%				
Epoch 19/100,	Train Loss: 1.4430,	Train Accuracy: 47.90%,	Test Loss: 1.4662,	Test Accurac
y: 47.38%				
Epoch 20/100,	Train Loss: 1.4056,	Train Accuracy: 50.20%,	Test Loss: 1.4032,	Test Accurac
y: 49.06%				
Epoch 21/100,	Train Loss: 1.3656,	Train Accuracy: 50.30%,	Test Loss: 1.4452,	Test Accurac
y: 46.39%				
Epoch 22/100,	Train Loss: 1.3847,	Train Accuracy: 49.40%,	Test Loss: 1.3860,	Test Accurac
y: 49.23%				
Epoch 23/100,	Train Loss: 1.3344,	Train Accuracy: 52.20%,	Test Loss: 1.3852,	Test Accurac
y: 48.55%				
Epoch 24/100,	Train Loss: 1.2814,	Train Accuracy: 53.80%,	Test Loss: 1.3503,	Test Accurac
y: 49.64%				
Epoch 25/100,	Train Loss: 1.2741,	Train Accuracy: 54.70%,	Test Loss: 1.3498,	Test Accurac
y: 50.46%				
Epoch 26/100,	Train Loss: 1.2313,	Train Accuracy: 56.60%,	Test Loss: 1.3201,	Test Accurac
y: 51.25%				
Epoch 27/100,	Train Loss: 1.2511,	Train Accuracy: 56.50%,	Test Loss: 1.3526,	Test Accurac
y: 50.99%				
Epoch 28/100,	Train Loss: 1.2025,	Train Accuracy: 57.20%,	Test Loss: 1.3112,	Test Accurac
y: 51.33%				
Epoch 29/100,	Train Loss: 1.1551,	Train Accuracy: 61.10%,	Test Loss: 1.2968,	Test Accurac
y: 52.56%				
Epoch 30/100,	Train Loss: 1.1251,	Train Accuracy: 60.50%,	Test Loss: 1.2858,	Test Accurac
y: 52.80%				
Epoch 31/100,	Train Loss: 1.1293,	Train Accuracy: 61.80%,	Test Loss: 1.2513,	Test Accurac
y: 53.94%				
Epoch 32/100,	Train Loss: 1.0751,	Train Accuracy: 64.60%,	Test Loss: 1.2418,	Test Accurac
y: 54.06%				
Epoch 33/100,	Train Loss: 1.0965,	Train Accuracy: 60.80%,	Test Loss: 1.2569,	Test Accurac
y: 53.34%				
Epoch 34/100,	Train Loss: 1.0706,	Train Accuracy: 63.30%,	Test Loss: 1.2326,	Test Accurac
y: 54.65%				
Epoch 35/100,	Train Loss: 1.0269,	Train Accuracy: 65.00%,	Test Loss: 1.2449,	Test Accurac
y: 54.14%				
Epoch 36/100,	Train Loss: 1.0039,	Train Accuracy: 65.30%,	Test Loss: 1.1951,	Test Accurac
y: 56.59%				
Epoch 37/100,	Train Loss: 0.9901,	Train Accuracy: 66.50%,	Test Loss: 1.2094,	Test Accurac
y: 56.02%				
Epoch 38/100,	Train Loss: 0.9855,	Train Accuracy: 67.40%,	Test Loss: 1.2605,	Test Accurac
y: 53.81%				
Epoch 39/100,	Train Loss: 0.9542,	Train Accuracy: 69.70%,	Test Loss: 1.2410,	Test Accurac
y: 54.56%				
Epoch 40/100,	Train Loss: 0.9256,	Train Accuracy: 68.20%,	Test Loss: 1.1748,	Test Accurac
y: 56.77%				
Epoch 41/100,	Train Loss: 0.8865,	Train Accuracy: 69.60%,	Test Loss: 1.2017,	Test Accurac
y: 56.30%				
Epoch 42/100,	Train Loss: 0.8975,	Train Accuracy: 70.40%,	Test Loss: 1.2456,	Test Accurac
y: 54.46%				
Epoch 43/100,	Train Loss: 0.8848,	Train Accuracy: 71.10%,	Test Loss: 1.1492,	Test Accurac
y: 58.74%				
Epoch 44/100,	Train Loss: 0.8544,	Train Accuracy: 70.70%,	Test Loss: 1.2504,	Test Accurac
y: 54.91%				
Epoch 45/100,	Train Loss: 0.8366,	Train Accuracy: 72.00%,	Test Loss: 1.1626,	Test Accurac
y: 57.88%				
Epoch 46/100,	Train Loss: 0.8217,	Train Accuracy: 72.90%,	Test Loss: 1.2054,	Test Accurac
y: 56.76%				
Epoch 47/100,	Train Loss: 0.8027,	Train Accuracy: 74.40%,	Test Loss: 1.1529,	Test Accurac
y: 58.75%				
Epoch 48/100,	Train Loss: 0.7676,	Train Accuracy: 75.20%,	Test Loss: 1.1577,	Test Accurac
y: 58.19%				
Epoch 49/100,	Train Loss: 0.7610,	Train Accuracy: 75.10%,	Test Loss: 1.1631,	Test Accurac
y: 57.99%				
Epoch 50/100,	Train Loss: 0.7847,	Train Accuracy: 74.20%,	Test Loss: 1.1579,	Test Accurac
y: 58.04%				
Epoch 51/100,	Train Loss: 0.7210,	Train Accuracy: 77.30%,	Test Loss: 1.2373,	Test Accurac
y: 55.98%				
Epoch 52/100,	Train Loss: 0.7384,	Train Accuracy: 75.70%,	Test Loss: 1.1618,	Test Accurac
y: 58.23%				
Epoch 53/100,	Train Loss: 0.7185,	Train Accuracy: 76.70%,	Test Loss: 1.1631,	Test Accurac
y: 58.61%				

Epoch 54/100,	Train Loss: 0.7062,	Train Accuracy: 77.40%,	Test Loss: 1.1857,	Test Accurac
y: 57.77%				
Epoch 55/100,	Train Loss: 0.6884,	Train Accuracy: 78.60%,	Test Loss: 1.1799,	Test Accurac
y: 58.56%				
Epoch 56/100,	Train Loss: 0.6882,	Train Accuracy: 77.90%,	Test Loss: 1.1330,	Test Accurac
y: 59.75%				
Epoch 57/100,	Train Loss: 0.6375,	Train Accuracy: 80.90%,	Test Loss: 1.2338,	Test Accurac
y: 56.90%				
Epoch 58/100,	Train Loss: 0.6419,	Train Accuracy: 80.90%,	Test Loss: 1.1585,	Test Accurac
y: 58.75%				
Epoch 59/100,	Train Loss: 0.6339,	Train Accuracy: 80.20%,	Test Loss: 1.1752,	Test Accurac
y: 58.66%				
Epoch 60/100,	Train Loss: 0.6133,	Train Accuracy: 80.90%,	Test Loss: 1.1830,	Test Accurac
y: 59.09%				
Epoch 61/100,	Train Loss: 0.5758,	Train Accuracy: 82.30%,	Test Loss: 1.1433,	Test Accurac
y: 60.14%				
Epoch 62/100,	Train Loss: 0.5664,	Train Accuracy: 82.60%,	Test Loss: 1.1998,	Test Accurac
y: 59.15%				
Epoch 63/100,	Train Loss: 0.5646,	Train Accuracy: 82.30%,	Test Loss: 1.1561,	Test Accurac
y: 59.26%				
Epoch 64/100,	Train Loss: 0.5597,	Train Accuracy: 83.10%,	Test Loss: 1.2595,	Test Accurac
y: 57.39%				
Epoch 65/100,	Train Loss: 0.5212,	Train Accuracy: 84.10%,	Test Loss: 1.1440,	Test Accurac
y: 60.21%				
Epoch 66/100,	Train Loss: 0.5071,	Train Accuracy: 85.20%,	Test Loss: 1.1739,	Test Accurac
y: 59.48%				
Epoch 67/100,	Train Loss: 0.4855,	Train Accuracy: 85.10%,	Test Loss: 1.1507,	Test Accurac
y: 60.36%				
Epoch 68/100,	Train Loss: 0.5062,	Train Accuracy: 84.20%,	Test Loss: 1.1972,	Test Accurac
y: 58.79%				
Epoch 69/100,	Train Loss: 0.4794,	Train Accuracy: 86.00%,	Test Loss: 1.1468,	Test Accurac
y: 60.39%				
Epoch 70/100,	Train Loss: 0.4641,	Train Accuracy: 87.40%,	Test Loss: 1.2114,	Test Accurac
y: 59.50%				
Epoch 71/100,	Train Loss: 0.4893,	Train Accuracy: 84.90%,	Test Loss: 1.2092,	Test Accurac
y: 59.16%				
Epoch 72/100,	Train Loss: 0.4466,	Train Accuracy: 86.20%,	Test Loss: 1.2072,	Test Accurac
y: 60.01%				
Epoch 73/100,	Train Loss: 0.4356,	Train Accuracy: 87.00%,	Test Loss: 1.2111,	Test Accurac
y: 59.34%				
Epoch 74/100,	Train Loss: 0.4133,	Train Accuracy: 88.30%,	Test Loss: 1.2133,	Test Accurac
y: 59.65%				
Epoch 75/100,	Train Loss: 0.4228,	Train Accuracy: 87.20%,	Test Loss: 1.2425,	Test Accurac
y: 59.41%				
Epoch 76/100,	Train Loss: 0.4151,	Train Accuracy: 88.60%,	Test Loss: 1.2763,	Test Accurac
y: 58.35%				
Epoch 77/100,	Train Loss: 0.4087,	Train Accuracy: 87.80%,	Test Loss: 1.1782,	Test Accurac
y: 61.12%				
Epoch 78/100,	Train Loss: 0.4053,	Train Accuracy: 87.50%,	Test Loss: 1.2756,	Test Accurac
y: 58.92%				
Epoch 79/100,	Train Loss: 0.3789,	Train Accuracy: 89.10%,	Test Loss: 1.2455,	Test Accurac
y: 59.45%				
Epoch 80/100,	Train Loss: 0.3968,	Train Accuracy: 88.50%,	Test Loss: 1.2935,	Test Accurac
y: 57.38%				
Epoch 81/100,	Train Loss: 0.3547,	Train Accuracy: 91.20%,	Test Loss: 1.2564,	Test Accurac
y: 60.38%				
Epoch 82/100,	Train Loss: 0.3620,	Train Accuracy: 88.90%,	Test Loss: 1.2921,	Test Accurac
y: 58.83%				
Epoch 83/100,	Train Loss: 0.3511,	Train Accuracy: 90.20%,	Test Loss: 1.2240,	Test Accurac
y: 60.12%				
Epoch 84/100,	Train Loss: 0.3478,	Train Accuracy: 89.60%,	Test Loss: 1.2884,	Test Accurac
y: 58.89%				
Epoch 85/100,	Train Loss: 0.2962,	Train Accuracy: 93.40%,	Test Loss: 1.3445,	Test Accurac
y: 58.66%				
Epoch 86/100,	Train Loss: 0.3305,	Train Accuracy: 90.90%,	Test Loss: 1.2106,	Test Accurac
y: 60.61%				
Epoch 87/100,	Train Loss: 0.3299,	Train Accuracy: 89.60%,	Test Loss: 1.2826,	Test Accurac
y: 59.09%				
Epoch 88/100,	Train Loss: 0.3205,	Train Accuracy: 90.70%,	Test Loss: 1.3072,	Test Accurac
y: 59.23%				
Epoch 89/100,	Train Loss: 0.2980,	Train Accuracy: 92.20%,	Test Loss: 1.3001,	Test Accurac
y: 60.14%				

Epoch 90/100, Train Loss: 0.2913, Train Accuracy: 91.40%, Test Loss: 1.2319, Test Accuracy: 60.81%

Epoch 91/100, Train Loss: 0.2936, Train Accuracy: 91.60%, Test Loss: 1.4124, Test Accuracy: 57.14%

Epoch 92/100, Train Loss: 0.2785, Train Accuracy: 92.20%, Test Loss: 1.2426, Test Accuracy: 60.88%

Epoch 93/100, Train Loss: 0.2576, Train Accuracy: 93.20%, Test Loss: 1.3336, Test Accuracy: 60.34%

Epoch 94/100, Train Loss: 0.2824, Train Accuracy: 92.60%, Test Loss: 1.3045, Test Accuracy: 59.06%

Epoch 95/100, Train Loss: 0.2692, Train Accuracy: 92.90%, Test Loss: 1.3033, Test Accuracy: 59.90%

Epoch 96/100, Train Loss: 0.2609, Train Accuracy: 92.40%, Test Loss: 1.2357, Test Accuracy: 61.12%

Epoch 97/100, Train Loss: 0.2752, Train Accuracy: 93.00%, Test Loss: 1.3697, Test Accuracy: 58.30%

Epoch 98/100, Train Loss: 0.2513, Train Accuracy: 93.20%, Test Loss: 1.2528, Test Accuracy: 61.26%

Epoch 99/100, Train Loss: 0.2147, Train Accuracy: 94.30%, Test Loss: 1.2809, Test Accuracy: 60.14%


Epoch 100/100, Train Loss: 0.2152, Train Accuracy: 95.20%, Test Loss: 1.2874, Test Accuracy: 60.59%

In [118]:

```
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(train_losses, label='Train Loss')
plt.plot(test_losses, label='Test Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Test Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(train_accuracies, label='Train Accuracy')
plt.plot(test_accuracies, label='Test Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy (%)')
plt.title('Training and Test Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```



In [119]:

```
def imshow(img):
    img = img / 2 + 0.5
    np_img = img.numpy()
    plt.imshow(np.transpose(np_img, (1, 2, 0)))
```



```
plt.axis('off')
plt.show()

def test_random_image(model, loader, device):
    model.eval()
    images, labels = next(iter(loader))
    images, labels = images.to(device), labels.to(device)

    import random
    index = random.randint(0, images.size(0) - 1)
    image = images[index].unsqueeze(0)
    label = labels[index].item()

    output = model(image)
    _, predicted = torch.max(output, 1)
    predicted = predicted.item()

    imshow(image.cpu().squeeze())
    print(f'Predicted: {predicted}, Actual: {label}')

test_random_image(model, test_loader, device)
```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



Predicted: 1, Actual: 1

In [ ]: